

Feast_NNI_BentoML_DB

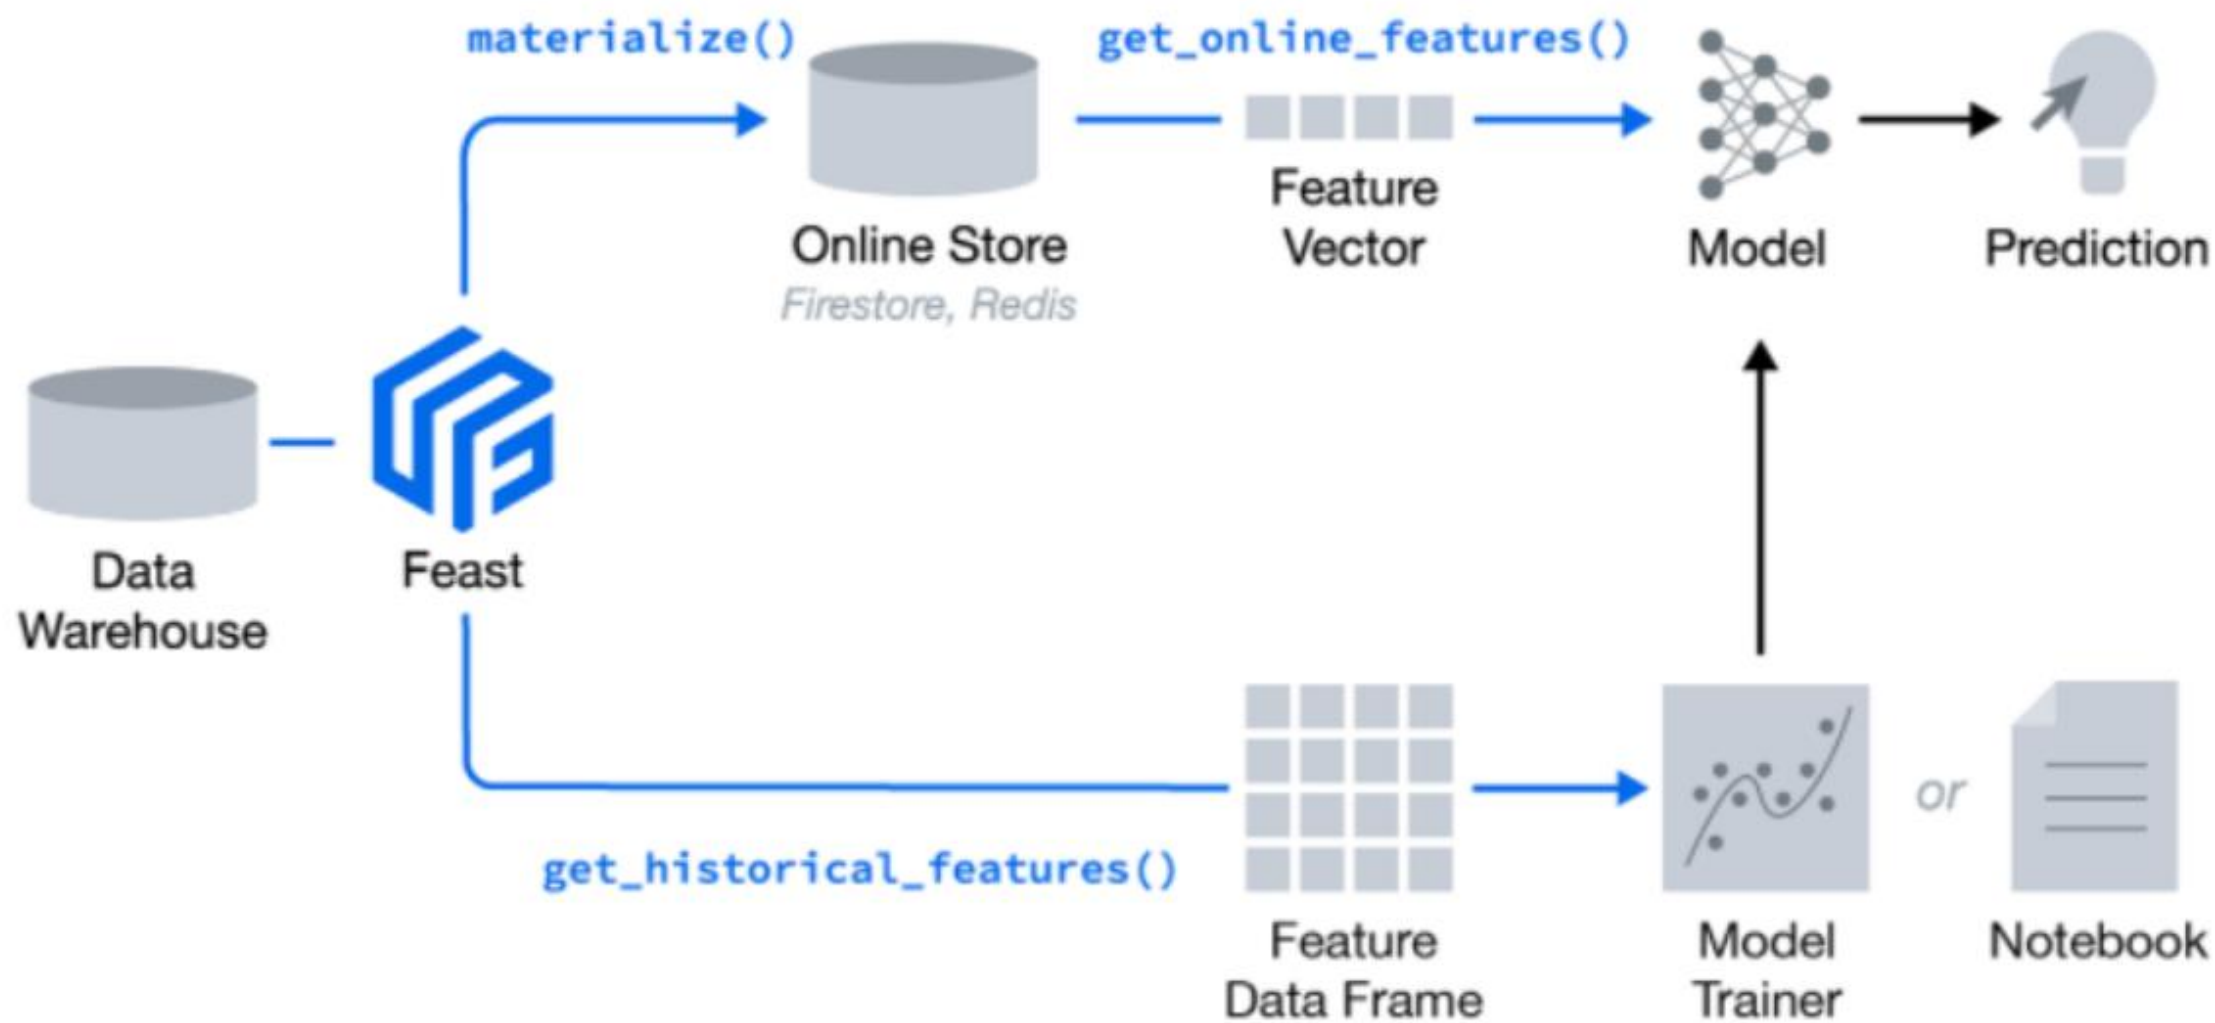
210901

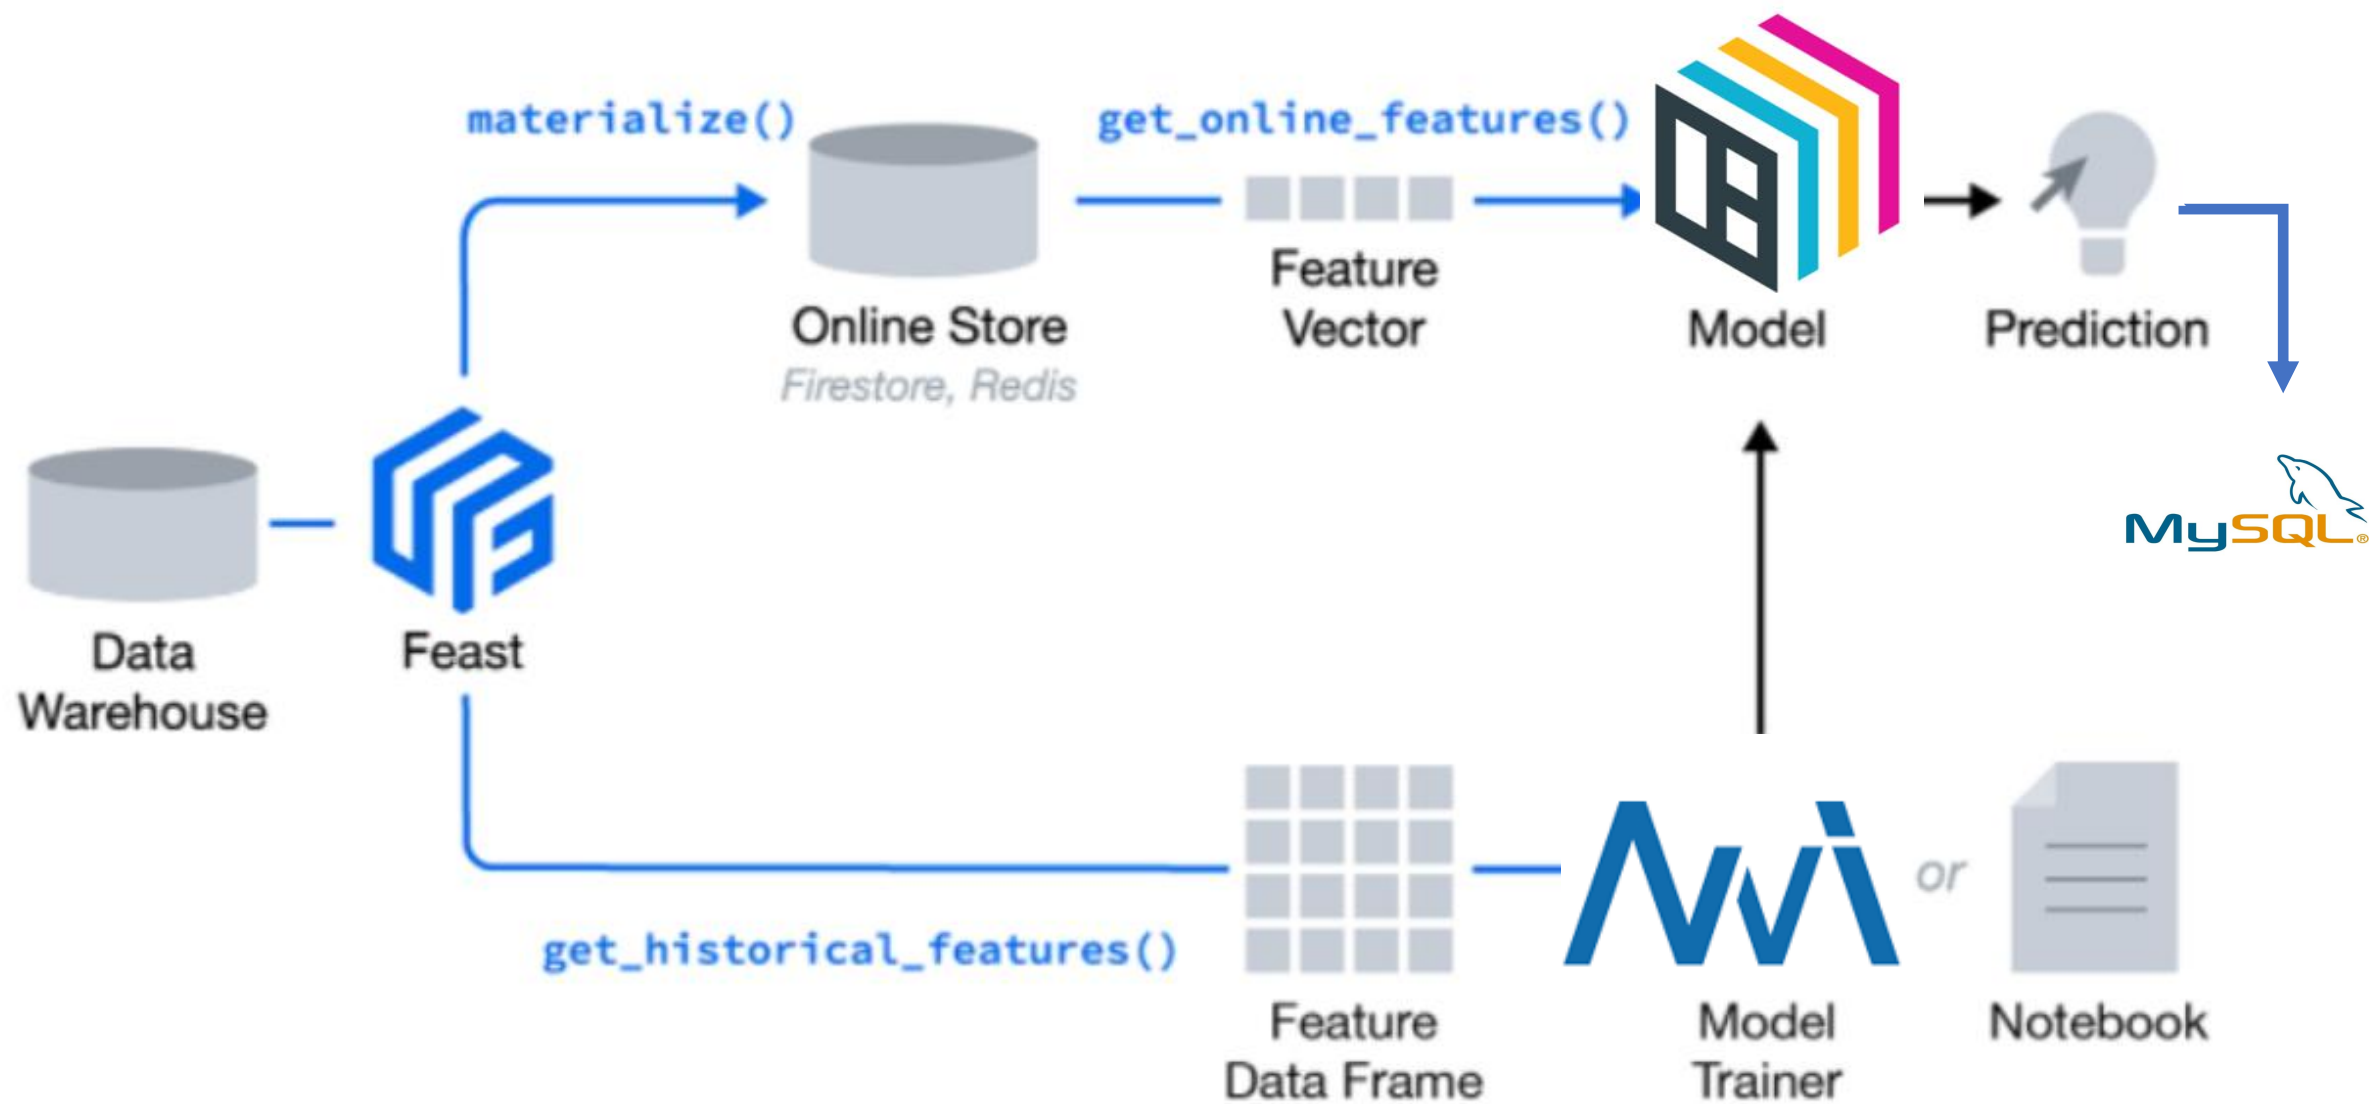
AI_Ops

이민준



FEAST











Parquet



Feast



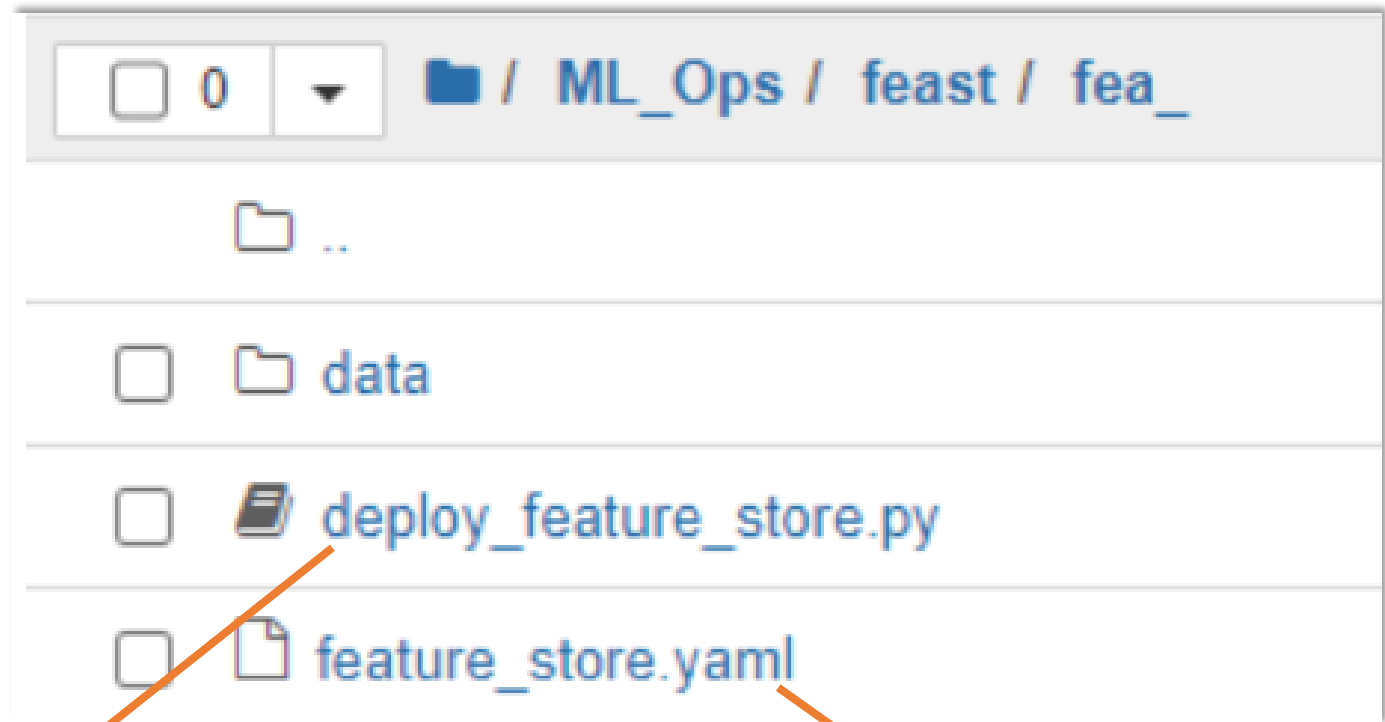


☐  ppr_data_.parquet

아직은 parquet만 지원

```
> feast init fea_
```

- 'fea_'



```
dr_lauren_stat = FileSource(  
    path="/workspace/ML_Ops/feast/fea_/data/ppr_data_.parquet",  
    event_timestamp_column="event_timestamp",
```

```
    dr_lauren_view = FeatureView(  
        Feature(name= columns_name, dtype = ~), x num_cols  
    )
```

```
project: fea_  
registry: data/registry.db  
provider: local  
online_store:  
    path: data/online_store.db
```

```
root:ML_Ops/feast/fea_
▶ feast apply
Registered entity ticket id
Registered feature view dr_lauren_stat
Deploying infrastructure for dr_lauren_stat
(base)
```



Feast_nni



```
### entity_df
parquet_ = pd.read_parquet(f'{path_}/data/ppr_data_.parquet', engine='pyarrow')
orders = parquet_[['ticket_id', 'event_timestamp']]

# Retrieve training data
training_df = fs.get_historical_features(
    entity_df=orders,
    features=[
        "dr_lauren_stat:time",
        "dr_lauren_stat:weekday",
        "dr_lauren_stat:weekend",
        "dr_lauren_stat:instlo_1",
        "dr_lauren_stat:instlo_2",
        "dr_lauren_stat:inst_code",
        "dr_lauren_stat:sysname_lo",
        "dr_lauren_stat:sysname_eq",
        "dr_lauren_stat:ntt_label",
    ],
).to_df()
```

	event_timestamp	ticket_id	time	weekday	weekend	instlo_1	instlo_2	inst_code	sysname_lo	sysname_eq	ntt_label
0	2021-01-01 03:17:00+00:00	3604002	3.283333	4	0	9	47	531	117	0	1
1	2021-01-01 03:17:00+00:00	3604002	3.283333	4	0	9	47	531	117	0	1
2	2021-01-01 03:20:00+00:00	3604035	3.333333	4	0	12	84	340	913	0	1
3	2021-01-01 03:20:00+00:00	3604035	3.333333	4	0	12	84	340	913	0	1
4	2021-01-02 08:38:00+00:00	3624904	8.633333	5	1	10	101	435	868	0	1
...
14730	2021-06-30 18:33:00+00:00	9346264	18.550000	2	0	2	67	400	173	0	0
14731	2021-06-30 18:35:00+00:00	9346310	18.583333	2	0	2	67	375	1506	0	0
14732	2021-06-30 18:35:00+00:00	9346329	18.583333	2	0	2	67	382	1978	0	0
14733	2021-06-30 18:40:00+00:00	9346522	18.666667	2	0	2	122	1504	161	0	0
14734	2021-06-30 18:40:00+00:00	9346522	18.666667	2	0	2	122	1504	161	0	0




Neural Network Intelligence



```
{  
  "model_n": {  
    "_type": "choice",  
    "_value": ["lgb", "xgb"]  
  },  
  "ne": {  
    "_type": "choice",  
    "_value": [50, 60, 70, 80, 90, 100]  
  },  
  "seed": {  
    "_type": "choice",  
    "_value": [42, 32, 16, 8]  
  }  
}
```

nmi



Experiment

ID: YpKXtrt9
gbm

Status
DONE

Best metric
0.982131

Start time
8/19/2021, 2:49:58 PM

End time
8/19/2021, 3:08:34 PM

Training platform
local

Tuner
TPE

Duration

Max duration 14399 min

18m 24s / 143998560 m

Trial numbers

Max trial No. 100

100 / 100

Concurrency 2

Running	Succeeded	Stopped	Failed	Waiting
0	100	0	0	0

Log directory
/root/nni-experiments/YpKXtrt9

Trial command
python test_py --gpu 0 --request-from-nni

Top trials


Max Min Display top 10

Trial No.	ID	Duration	Status	Default metric
28	ZI4c5	6s	SUCCEEDED	0.982131
16	fNG2d	5s	SUCCEEDED	0.982131
53	hRRWH	5s	SUCCEEDED	0.981226
45	IoTdH	6s	SUCCEEDED	0.981226
61	oM6ir	6s	SUCCEEDED	0.981
35	EH8j3	6s	SUCCEEDED	0.981
24	lpuQa	6s	SUCCEEDED	0.981
19	LIAR6	6s	SUCCEEDED	0.981
66	PN074	6s	SUCCEEDED	0.980321
38	CnsuM	6s	SUCCEEDED	0.980321

<input type="checkbox"/>	0	▼	/ ML_Ops / feast / models_
			..
<input type="checkbox"/>			lgb_acc_0.95273_auc_0.9925_.pth.tar
<input type="checkbox"/>			lgb_acc_0.95454_auc_0.99256_.pth.tar
<input type="checkbox"/>			lgb_acc_0.95589_auc_0.99372_.pth.tar
<input type="checkbox"/>			lgb_acc_0.95634_auc_0.99409_.pth.tar
<input type="checkbox"/>			lgb_acc_0.95838_auc_0.99412_.pth.tar
<input type="checkbox"/>			lgb_acc_0.96019_auc_0.99451_.pth.tar
<input type="checkbox"/>			lgb_acc_0.96042_auc_0.9949_.pth.tar



BENTOML



Pack.py

```
@env(infer_pip_packages=True)
@artifacts([SklearnModelArtifact('model')])
class Dr_lauren_classifier(BentoService):
    """
    A minimum prediction service exposing a Scikit-learn model
    """

    @api(input=DataframeInput(), batch=True)
    def predict(self, df: pd.DataFrame):
        """
        An inference API named `predict` with Dataframe input adapter, which codifies
        how HTTP requests or CSV files are converted to a pandas Dataframe object as the
        inference API function input
        """
        return self.artifacts.model.predict(df)
```

Package_save.py

```
# create dr_lauren_service instance
dr_lauren_service = pack_.Dr_lauren_classifier()

# get_model
mo_ = torch.load('/workspace/ML_Ops/feast/fea_/models_/xgb_acc_0.97829_auc_0.99705_.pth.tar')['model']

# dr_lauren_classifier와 'model'로 패키징됨
dr_lauren_service.pack('model', mo_)

# 경로 저장
saved_path = dr_lauren_service.save()
```

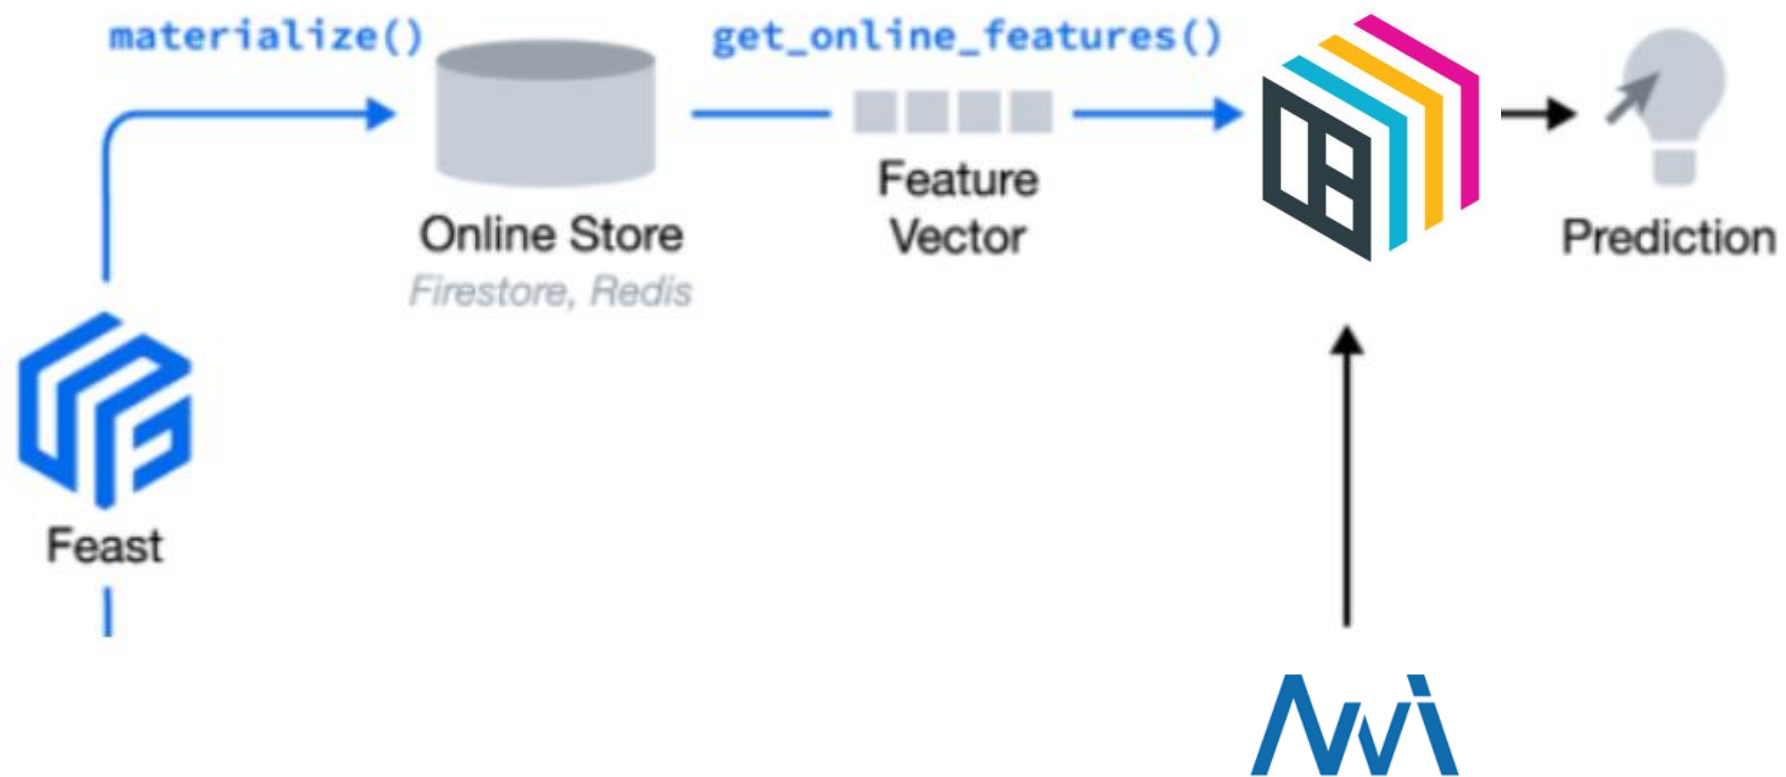
```
> Python Package_save.py
```

```
root:ML_Ops/feast/fea_
▶ ls /root/bentoml/repository/Dr_lauren_classifier
20210819063311_8F9A48  20210819074959_34A218  20210819075428_1FF3A7  20210819080455_1C8532
(base)
root:ML_Ops/feast/fea_
▶ ls /root/bentoml/repository/Dr_lauren_classifier/20210819080455_1C8532
bentoml-init.sh  docker-entrypoint.sh  docs.json  environment.yml  python_version  requirements.txt
bentoml.yml      Dockerfile             Dr_lauren_classifier  MANIFEST.in     README.md       setup.py
(base)
```



```
> Bentoml serving Dr_lauren_classifier:latest
```

```
root:/workspace/ML_Ops/feast
▶ bentoml serve Dr_lauren_classifier:latest
[2021-08-22 08:48:12,763] INFO - Getting latest version Dr_lauren_classifier:20210819080455_1C8532
[2021-08-22 08:48:12,775] INFO - Starting BentoML API proxy in development mode..
[2021-08-22 08:48:12,778] INFO - Starting BentoML API server in development mode..
[2021-08-22 08:48:12,975] INFO - Micro batch enabled for API `predict` max-latency: 20000 max-batch-size 4000
[2021-08-22 08:48:12,975] INFO - Your system nofile limit is 1048576, which means each instance of microbatch service is able to hold this number of connections at same time. You can increase the number of file descriptors for the server process, or launch more microbatch instances to accept more concurrent connection.
===== Running on http://0.0.0.0:5000 =====
(Press CTRL+C to quit)
* Serving Flask app "Dr_lauren_classifier" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://127.0.0.1:51222/ (Press CTRL+C to quit)
```






```
> Feast materialize 2021-01-01T00:00:00 2021-08-01T00:00:00
```

Feast_bentoml

```
► feast materialize 2021-01-01T00:00:00 2021-08-01T00:00:00  
Materializing 1 feature views from 2021-01-01 00:00:00+00:00 to 2021-08-01 00:00:00+00:00 into the sqlite online store.  
  
dr_lauren_stat:  
/opt/conda/lib/python3.8/site-packages/pyarrow/pandas_compat.py:1027: DeprecationWarning: `np.float` is a deprecated alias  
for the builtin `float`. To silence this warning, use `float` by itself. Doing this will not modify any behavior and is saf  
e. If you specifically wanted the numpy scalar type, use `np.float64` here.  
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations  
  'floating': np.float,  
100%|██████████████████████████████████████████████████████████████████████████████| 10784/10784 [00:03<00:00, 2814.03it/s]  
(base)
```

ticket_id	event_timestamp
9360646	2021-07-01 02:47:00
9361358	2021-07-01 03:11:00
9363518	2021-07-01 04:25:00
9365171	2021-07-01 05:29:00
9365176	2021-07-01 05:29:00
...	...
9995149	2021-07-15 18:37:00
9995205	2021-07-15 18:39:00
9995206	2021-07-15 18:39:00
9995241	2021-07-15 18:40:00
9995586	2021-07-15 18:49:00

```
def get_entities(path_, from_time_):  
    parquet_ = pd.read_parquet(path_, engine='pyarrow')  
    orders = parquet_[['ticket_id', 'event_timestamp']]  
  
    ### 뽐을 rows의 entity == key 라고 볼 수 있음  
    new_orders = orders[orders['event_timestamp'] >= from_time_]  
    new_orders.drop_duplicates(ignore_index=True, inplace=True)  
    return new_orders
```


Test_entity

```
online_ = fs_.get_online_features(  
    entity_rows=[{"ticket_id": i} for i in entity_df['ticket_id']],  
    features=[  
        "dr_lauren_stat:time",  
        "dr_lauren_stat:weekday",  
        "dr_lauren_stat:weekend",  
        "dr_lauren_stat:instlo_1",  
        "dr_lauren_stat:instlo_2",  
        "dr_lauren_stat:inst_code",  
        "dr_lauren_stat:sysname_lo",  
        "dr_lauren_stat:sysname_eq",  
        "dr_lauren_stat:ntt_label",  
    ],  
)
```

	time	weekday	weekend	instlo_1	instlo_2	inst_code	sysname_lo	sysname_eq	ntt_label
0	2.783333	3	0	3	15	781	1584	0	1
1	3.183333	3	0	12	38	693	972	0	1
2	4.416667	3	0	3	53	293	2200	0	1
3	5.483333	3	0	11	37	91	834	0	1
4	5.483333	3	0	11	37	525	820	0	1
...
1368	18.616667	3	0	2	67	400	173	0	0
1369	18.650000	3	0	2	67	375	1506	0	0
1370	18.650000	3	0	2	67	382	1978	0	0
1371	18.666667	3	0	9	65	642	2061	0	1
1372	18.816667	3	0	14	128	1215	2047	0	1

Request

```
def get_from_bentoml(online_df_):  
    response = requests.post("http://127.0.0.1:5000/predict", data=online_df_.iloc[:, :-1].to_json())  
    infer_ = response.text  
    predict_ = infer_[1:-1].split(', ')  
    return predict_  
  
def get_infer_df_(entity_df_, online_df_):  
    infer_df = entity_df_.reset_index(drop=True)  
    infer_df['label'] = online_df_['ntt_label']  
    infer_df['predict'] = get_from_bentoml(online_df_)  
    return infer_df
```



Feast_bentoml

[illegible]

To_DB



data create, insert, update

```
def get_commit_sql(conn_, sql_):  
    with conn_.cursor() as cursor:  
        cursor.execute(sql_)  
        result = cursor.fetchall()  
        conn_.commit()
```


data read

```
def get_data_sql(conn_, sql_):  
    with conn_.cursor() as cursor:  
        cursor.execute(sql_)  
        result = cursor.fetchall()  
    return result
```

```
def get_columns_(df_):  
    columns_str = ', '.join([i for i in infer_df.columns])  
    return columns_str
```

```
def get_values_(df_):  
    value_ = []  
    for row_ in df_.values:  
        value_.append("(" + ', '.join("'" + str(i) + "'" for i in row_) + ")")  
  
    values_str = ', '.join(value_)  
    return values_str
```



```
mysql> select * from infer_df limit 30;
```

ticket_id	event_timestamp	label	predict
9360646	2021-07-01 02:47:00	1	1
9361358	2021-07-01 03:11:00	1	1
9363518	2021-07-01 04:25:00	1	1
9365171	2021-07-01 05:29:00	1	1
9365176	2021-07-01 05:29:00	1	1
9368107	2021-07-01 07:27:00	1	1
9370665	2021-07-01 08:58:00	1	1
9370880	2021-07-01 09:07:00	1	1
9370921	2021-07-01 09:07:00	1	1
9371927	2021-07-01 09:41:00	1	1
9371998	2021-07-01 09:45:00	0	1
9372121	2021-07-01 09:48:00	0	1
9372211	2021-07-01 09:51:00	1	1
9372328	2021-07-01 09:51:00	1	1
9372338	2021-07-01 09:54:00	0	1
9372362	2021-07-01 09:54:00	0	1
9373363	2021-07-01 10:25:00	1	1
9373412	2021-07-01 10:27:00	1	1
9373575	2021-07-01 10:33:00	1	1
9373598	2021-07-01 10:32:00	1	1
9376262	2021-07-01 12:01:00	1	1
9379021	2021-07-01 13:30:00	1	1
9381033	2021-07-01 14:31:00	1	1
9381165	2021-07-01 14:33:00	1	1
9384423	2021-07-01 16:11:00	1	1
9384623	2021-07-01 16:18:00	1	1
9384676	2021-07-01 16:18:00	1	1
9384712	2021-07-01 16:21:00	1	1
9384796	2021-07-01 16:22:00	1	1
9384892	2021-07-01 16:26:00	1	1

```
30 rows in set (0.00 sec)
```


느낀 점

감사합니다.